

$$\frac{188}{240} = 78.3\%$$

Assignment 1

Exploratory Data Analysis and the Implementation of K-Nearest Neighbor and Classification Tree Models on Dry Bean Data

A.D. van der Merwe
Department of Computer Science
University of Stellenbosch
24923273
24923273@sun.ac.za

W198: This report explores...

Abstract—In this report, the application of classification algorithms is explored on the Dry Bean dataset, with a focus on data analysis and data preparation. The primary goal is to prepare the data and then construct two classification models, k-nearest neighbours (KNN) and decision trees on the Dry Bean dataset and evaluate their performance on predictions of bean types. The characteristics of each feature and visualisation is explored to determine if there are any data quality issues. The data is prepared for use in the classification algorithms, which includes standardisation for KNN and post-pruning for the decision tree. K-fold cross-validation assesses the effectiveness of the model and if the model is deterministic, with performance measured by accuracy and F1-score. The report shows that both models can be tailored to handle data quality issues and highlights the strength of each approach when beans are classified.

I. INTRODUCTION

Exploratory data analysis EDA is crucial for machine learning as it provides knowledge about the dataset and how to approach the task at hand. More knowledge of a dataset can help make a decision on which modeling techniques to apply to the dataset.

W123 The approach began with a thorough characterisation of the Dry Bean dataset, to better understand the nature and distributions of the features by the calculation of descriptive statistics. Data quality issues are identified such as missing values, class imbalances and outliers. Data preparation techniques are applied to the data data quality issues to ensure optimal performance of the classification algorithms that are implemented. The classification algorithms implemented are a k-nearest neighbours (KNN) and decision tree classification algorithm. For KNN, feature standardisation and parameter tuning is conducted to enhance the performance of the classification model, while, for the decision tree classification model, tree induction and post-pruning techniques are applied to prevent overfitting. To evaluate these models k-fold cross-validation is applied and different performance features like accuracy and F1-score is used. This comprehensive approach aims to ensure a good understanding of the dataset, analyse the

data, address any data quality issues, and determine the most effective classification model for the Dry Bean dataset.

The rest of this paper is structured as follows: Section II provides background information on data quality issues, data preparation techniques, KNN classification algorithms and decision tree classification algorithms. Section III provides the methodology used to analyse and prepare the Dry Bean dataset and to build a KNN and decision tree classification algorithm on the data. The empirical procedure is given in Section IV and the research results is given in Section V.

II. BACKGROUND

This section presents background information on EDA and techniques used to clean and transform data for analysis or to construct models. Additionally, information on the C4.5 decision tree and k-nearest neighbours classification models.

A. Exploratory Data Analysis

Exploratory data analysis EDA first introduced by John Tukey in 1977 [8]. Tukey emphasised that there are three key strategies in data analysis, namely:

- Graphical presentation: the use of visual tools, such as plots, to explore, analyse and understand data.
- Flexibility in viewpoint and facilities: encourages creative and diverse approaches to analyse data that allows for dynamic perspective and adaptive tools.
- Search for parsimony and simplicity: To find simple and clear explanations in the data while unnecessary complexity is avoided.

1) Data Quality Issues:

a) *Missing values*: The absence of data in a dataset where a value should be present is known as a missing value. There are three types of missing values, namely missing at random (MAR), missing completely at random (MCAR) and nonignorable (NI) missing values [1].

MCAR refers to a situation where missing values are distributed randomly throughout the dataset, independent of both observed and unobserved data. Under MCAR, modern

methods such as missing value imputation can be used and still produce unbiased estimates, though some uncertainty is introduced. When modern methods are used to handel MCAR missing values, the statistical power reduces compared to complete data.

A missing value is MAR if the likelihood of the missing data on the variable is unrelated to the value of that variable itself, once other variables are taken into account. These other features help explain why data may be missing and is known as mechanisms. Some common mechanisms include factors like education, income or age. The MAR assumption holds if the pattern of missingness is random when explained by the mechanisms.

If data is missing in a systematic way and can neither be classified as MAR or MCAR, then is the missing value classified as NI. To model this type of data is complex.

There are three ways to deal with these missing values. The first option explored is to use machine learning algorithms that is robust to missing values. The second option is to remove the instances that contains the missing values. This option does run the risk of valuable information that may be lossed. The third option is to impute the missing values.

When a missing value is imputed from a categorical feature, the most frequently occurring value, known as the mode, is commonly used. When a numerical feature is imputed, the median is typically used when the data contains outliers and the mean is used when the data contains no outliers.

b) Outliers: When one or more variables, of an instance in the dataset, have values that significantly differ from the overall distribution of the data, it is known as an outlier. There are two types of outliers, namely invalid outliers that is data points included in the dataset due to errors, and valid outliers that is correct values that are legitimately different from the rest of the data.

There are three techniques to cope with outliers. The first technique is to remove the outliers using statistical techniques. The second technique is to keep the outliers and apply robust estimators that are less sensitive to the influence of the outliers. Alternatively, clipping can be performed to the outliers to limit the range of the outliers or to impute a value to replace the outlier with. The third technique is to remove the outliers directly within the machine learning process by the use of algorithms designed to be robust to outliers.

One approach to perform outlier detection is by the use of a statistical method known as the Interquartile Range (IQR) method. The equation to calculated the IQR of an input feature \mathbf{p} , is as follows:

$$IQR^{\mathbf{p}} = Q_3^{\mathbf{p}} - Q_1^{\mathbf{p}} \quad (1)$$

where $Q_1^{\mathbf{p}}$ and $Q_3^{\mathbf{p}}$ is the first and third quartile of the input feature \mathbf{p} respectively. The equations to determine if an instance of \mathbf{p} is an outlier, is as follows:

$$\mathbf{p}_i < Q_1^{\mathbf{p}} - 1.5 \times IQR^{\mathbf{p}} \quad (2)$$

or if:

$$\mathbf{p}_i > Q_3^{\mathbf{p}} + 1.5 \times IQR^{\mathbf{p}} \quad (3)$$

where \mathbf{p}_i is the i -th instance of the input feature \mathbf{p} .

Another technique to apply outlier detection is by using visual methods such as a boxplot, that displays the distribution of data based on quartiles and highlights potential outliers, or by using a scatter plot to visualise the the relationship between two features and can reveal outliers as points that deviate significantly from the overall pattern.

c) Noise: Noise refers to random or irrelevant variations in the data that can obscure patterns. There is two types of noise. The first typeof noise is stochastic noise, with random variation in data values, a zero mean in the variation and small deviations. If the deviation of stochastic noise is large, the feature becomes irrelevant. The second type of noise is systematic noise, where all values shifted in a particular direction, such as due to poorly calibrated instruments.

When handling noise, systematic noise can be corrected by use of a few different techniques. One technique is to make use of standardisation. Stochastic noise is handled by the machine learning process, but caution should be used as if noise is present and the model is too complex, the model will overfit.

There are a few techniques used to remove noise in the dataset, namely edited nearest neighbour (ENN), repeated edited nearest neighbour (RENN), all k-nearest neighbour (ANN), cut edges weight statistic (CEWS) and Tomek links. A Tomek link is a pair of data instances \mathbf{a} and \mathbf{b} where there exists a third instance \mathbf{c} such that:

$$d(\mathbf{a}, \mathbf{c}) \leq d(\mathbf{a}, \mathbf{b}) \quad (4)$$

or

$$d(\mathbf{b}, \mathbf{c}) \leq d(\mathbf{a}, \mathbf{b}) \quad (5)$$

where d is the distance between the two instances.

If \mathbf{a} and \mathbf{b} forma Tomek link, then either one of them is noise or both are borderline cases. In such cases both \mathbf{a} and \mathbf{b} is removed to improve the quality of the dataset.

d) Imbalanced class data: A dataset is refered to as imbalanced, if the target feature has significantly unequal representation among its different categories. A few approaches to handel imbalanced classes are, do to nothing, but this is not a great solution, as the train and test accuracy of the majority class may be very good and the train and test accuracy of the minority class will be bad. Another approach to handel imbalanced classes is to balance the training set by either the use of synthetic minority oversampling technique (SMOTE) to oversample the minority class or the use of Tomek links to undersampl the majority class.

By using Tomek links as described in Equation (4) and Equation (5), it is possible to obtain Tomek links between classes and to then remove the majority class from each Tomek link.

SMOTE is used to oversample the minority class. The process of SMOTE is to firstly ignore the majority class. The next step is to choose the KNN for each minority instance. SMOTE then produces new instances halfway between the the instance and the neighbours.

2) Data Preparation:

a) *Data type transformations*: Different machine learning algorithms have varying requirements for the types of input data. This necessitates the conversion of the data from one type to another.

Discretisation is the process of converting continuous features to discrete features by the creation of bins that group the continuous values into distinct categories. Each bin represents a specific range derived from the the original continuous features. There is a general trade-off when a decision has to be made about the number of bins to use. This trade off is that valuable information may be lossed, if the number of bins is very low and if the number of bins is high, there might be very few instances in each bin or a bin could end up being empty.

There are typically two methods to encode discrete feature as a numerical Feature. The first method is known as ordinal encoding. Ordinal encoding transforms discrete features that typically has a natural order into continuous features. An example of this is that if the unique values of a discrete feature is cold, mild and high will be encoded as 0, 1 and 2 respectively.

One hot encoding is typically used when there is no natural order of the unique values of the discrete feature. If the unique values are cloudy, sunny and rainy, the values will be encoded as 1 0 0 for cloudy, 0 1 0 for sunny and 0 0 1 for rainy.

b) *Feature selection*: The goal of any feature selection method is to find the smallest subset of informative features that preserves the overall performance of the model. A smaller subset of features is beneficial as the models fitted to the data will be simpler. The smaller set of features will also produce shorter train times and produce less overfitted models, therefore improving the generalisation of a model.

If the subset of features contains too few features, the model may underfit due to the lack of discriminant power to differentiate between classes. An underfitted model leads to more false positives and false negatives. Too many features in the new subset of features can introduce noise in the training data, causing the model to overfit and increase computational complexity.

There are four types of different descriptive features to distinguish from, namely:

- **Predictive** descriptive feature that provides useful information.
- **Interacting** descriptive feature is not informative about the target feature on its own, but can provide valuable information when combined with one or more other features.
- **Redundant** descriptive feature that has a strong correlation with another feature.
- **Irrelevant** descriptive feature that does not provide any useful information when a value is estimated for the target feature.

Ideally, after a feature selection approach is applied, will the subset of features include the predictive and interacting features while the irrelevant and redundant features are excluded.

There are three methods to perform subset selection on a dataset, namely filter methods, wrapper methods and embedded methods. Filter methods performs statistical tests between input feature and target feature to identify a subset of features that are more relevant than others. The statistical test to apply types of input features to a specific type of target feature is represented by Table I.

TABLE I: Statistical Tests for Evaluating Different Types of Input Features Against a Specific Type of Target Feature

Feature Type	Target Feature Type	
	Continuous	Categorical
Continuous	Pearson Correlation	LDA
Categorical	Anova	Chi-Square

Two common approaches to apply the wrapper feature selection technique is by the use of forward and backward feature selection that produces a best subset of features to train the model on. In forward and backward feature selection a feature is added or removed to the best subset of features respectively and then the performance of the model is compared when using the best subset of features and the best subset of features that has been modified.

Embedded subset selection techniques is when an algorithm has a build in feature to perform subset selection. Models like decision trees prune away splits on features that has little to no predictive power and techniques like regularisation is applied to the objective function of a model to reduce the predictive power of features that has little to no influence when predicting a target feature.

c) *Normalisation*: Normalisation or scaling of input features is required for some machine learning algorithms. Normalisation is typically required if the ranges of values for different input features differ in order of magnitudes and where the algorithm makes use of a distance based metric to generate classification models. Normalisation is also used when the minimum and maximum values are not known, to reduce the effect of outliers or to transform a range of values to a different range of values. One method of normalisation typically used is the Z-score normalisation also known as standardisation. The equation to perform standardisation is as follows:

$$Z_{i,j} = \frac{x_{i,j} - \mu_j}{\sigma_j} \quad (6)$$

where $x_{i,j}$ is the data value of the j -th feature for the i -th observation, μ_j and σ_j is the mean and standard deviation of the j -th feature across all observations respectively and $Z_{i,j}$ is the Z-score of the data value at the i -th observation and the j -th feature.

Standardisation transform all of the values of each feature to have a zero-mean and a unit variance.

B. K-Nearest Neighbours

The KNN algorithm was first proposed by Evelyn Fix and Joseph Hodges in 1951 [5] in a technical report that was never published. The report contained pioneering work on

nonparametric discriminant analysis and probability density estimation and laid the foundational principles for the KNN algorithm.

KNN is a type of instance-based learning that does not build a model based on the training dataset. Instead, KNN computes the distance between the instances in the test set and all of the instances in the training set to identify the nearest neighbours for each entry in the test set. KNN relies heavily on the distance metric in order to identify the instances in the training set closest to each of the test instances. The distance metric plays a crucial role in order to determine the similarity between data points, that directly influences the performance and accuracy of the algorithm. The most common distance metrics used in the KNN algorithm is the:

- Euclidean distance metric
- Manhattan distance metric
- Minkowski distance metric

The Euclidean distance metric equation is defined as follows:

$$d(\mathbf{x}_1, \mathbf{x}_2) = \sqrt{\sum_{n=1}^N (x_{1n} - x_{2n})^2} \quad (7)$$

The Manhattan distance metric equation is defined as follows:

$$d(\mathbf{x}_1, \mathbf{x}_2) = \sum_{n=1}^N |x_{1n} - x_{2n}| \quad (8)$$

The Minkowski distance metric equation is defined as follows:

$$d(\mathbf{x}_1, \mathbf{x}_2) = \left(\sum_{n=1}^N |x_{1n} - x_{2n}|^p \right)^{\frac{1}{p}} \quad (9)$$

where \mathbf{x}_1 and \mathbf{x}_2 is a vector containing all the of the features of two distinct instances in a dataset, $d(\mathbf{x}_1, \mathbf{x}_2)$ is the distance between the two vectors \mathbf{x}_1 and \mathbf{x}_2 , N denotes the dimension of the vectors \mathbf{x}_1 and \mathbf{x}_2 or equivalently the total number of features and x_{1n} and x_{2n} is the value of the n -th feature of the two distinct feature vectors \mathbf{x}_1 and \mathbf{x}_2 respectively.

After the distances between the instance from the test set is calculated against all of the instances in the training set, a distance vector \mathbf{d} is obtained. The distance vector is then sorted from smallest to largest distances and the first k distances is chosen from \mathbf{d} . The process to classify an unseen instance, \mathbf{x} , on k of the instances of the training set D is described by Algorithm 1.

Algorithm 1 k-Nearest Neighbours (kNN)

```

1: function KNN( $D, \mathbf{x}, k$ )
2:   for all  $\mathbf{x}_i \in D$  do
3:      $\mathbf{d} = \text{DISTANCE}(\mathbf{x}_i, \mathbf{x})$ 
4:   end for
5:    $\text{SORT}(\mathbf{d})$ 
6:    $S = \text{set of } k \text{ patterns in } D \text{ closest to } \mathbf{x}$ 
7:   Return class as majority class in  $S$ 
8: end function

```

Data quality
Issues vs
kNN?

C. C4.5 Decision Trees

The first implementation of a decision tree algorithm was introduced by Leo Brieman *et al* in 1984 and is known as classification and regression trees (CART) [3]. In 1986, Ross Quinlan implemented the iterative dichotomiser 3 (ID3) decision tree algorithm, which is able to cope with noise and missing values in the data [6]. In 1993, Ross Quinlan introduced the C4.5 decision tree algorithm, which is able to handle continuous features, imbalanced classes and introduced post-pruning [7].

- C4.5 generates a classification tree model that consists of:
- Leaf nodes that indicates different discrete classes
 - Decision nodes that represent tests on a feature and leads to further branches in the tree.

C4.5 is induced to overfit the training data, thus the tree classifier completely separates the classes of the training data and leads to poor generalisation on unseen data as well as outliers and noise will be in the leaves of the tree.

C4.5 uses entropy as a measure of impurity in a dataset with respect to the target feature. The entropy of a dataset will be 0, if all of the classes in the dataset has the same labels. If the frequency of the classes are all equal, then will the entropy of the dataset be equal to 1. The test on an input feature that minimises the impurity of a dataset, or equivalently maximises the information gain, is selected as the best possible feature to split on. The equation to calculate the entropy of a dataset is as follows:

$$H(D) = - \sum_{i=1}^M p(y_m) \log_M(p(y_m)) \quad (10)$$

and

$$p(y_m) = \frac{\text{freq}(y_m, D)}{|D|} \quad (11)$$

where D represents the dataset, y_m is the m -th class, M is the total number of distinct classes in the dataset, $p(y_m)$ is the probability of a class y_m occurring in D and where $\text{freq}(Y_m, D)$ is the number of times class y_m occurs in D .

The equation used to calculate the entropy due to the split, follows as:

$$H_x(D) = \sum_{o=1}^O p_o H(D_o) \quad (12)$$

and

$$p_o = \frac{|D_o|}{|D|} \quad (13)$$

where x is the input variable used to split D , O is the number of outcomes for x , $H(D_o)$ is the conditional entropy of dataset D_o with reference to the class distribution and where p_o is the probability of outcome o .

The equation used to calculate the information gain if D is partitioned on x , follows as:

$$\text{gain}(x) = H(D) - H_x(D) \quad (14)$$

The information gain calculation favours input features with many outcomes and introduces bias as the tree creates multiple branches, one for each unique value. The solution to the bias

is to normalise the information gain with entropy with respect to test outcomes. The equation for the gain ratio is as follows:

$$\text{gainRatio}(x) = (1 - F) \frac{\text{gain}(x)}{\text{splitInfo}(x)} \quad (15)$$

and

$$\text{splitInfo}(x) = - \sum_{o=1}^O p_o \log_O(p_o) \quad (16)$$

where F is the fraction of patterns in D for which the value of x is missing. The objective of the tree is to select x that maximise the gain ratio.

In order to handel numerical continuous features, C4.5 first sorts the number of unique values, denoted by I , found in the numeric input feature x to create the sequence, (v_1, v_2, \dots, v_I) ordered from smallest to largest. C4.5 performs $I-1$ splits between v_i and v_{i+1} and for each of these splits the information gain is calculated by use of Equation (15). The split with the highest calculated information gain will be chosen and the threshold to split the tree on is calculated as follows:

$$\epsilon = \frac{v_i + v_{i+1}}{2} \quad (17)$$

where ϵ is the threshold to split the tree on. C4.5 then generates two decision nodes with test $x < \epsilon$ and $x \geq \epsilon$.

C4.5 partitions the dataset D in the absence of missing values as follows:

$$w_j(D_o) = 1, w_j(D_q) = 0, \forall q \neq o \quad (18)$$

where each pattern j in D belongs to subset D_o with weight $w_j(D_o)$, provided that pattern j no missing value for the input feature x . If pattern j has a missing value for feature x , it is assigned to all outcomes with a weight given by

$$w_j(D_o) = w_j(D) p_o \quad (19)$$

where the probability of the outcome is now calculated as follows:

$$p_o = \frac{\sum_{x_j \in D \text{ with outcome } o} w_j(D)}{\sum_{x_j \in D \text{ with any known outcome}} w_j(D)} \quad (20)$$

To determine the optimal input feature x to split the data on, the process follows these steps:

- Calculate the entropy of the current dataset, $H(D)$ by use of Equation (10)
- Compute the conditional entropy for each input feature x , $H_x(D)$ by use of Equation (12)
- Calculate the information gain for each feature if D would be split on x by use of Equation (14)
- Using the calculated information gain, calculate the gain ratio by use of Equation (15)
- Select the feature with the highest gain ratio and split D into O subsets
- Repeat the process on each subset until the tree has been induced to completely separate the classes

Algorithm 2 describes the process to fully induce the decision tree.

Algorithm 2 Classification Tree Induction

```

1: function INDUCETREE( $D$ )
2:   if  $|D| = 0$  then
3:     Return leaf with default class
4:   end if
5:   if  $|D| > 0$  and  $\forall \mathbf{x} \in D$  the class is the same, i.e.  $y_m$ 
     then
6:     Return leaf with class label  $y_m$ , containing  $D$ 
7:   end if
8:   Select a test based on a single input variable
9:   Split  $D$  into  $D_1, D_2, \dots, D_o$ , where  $O$  is the number of
     outcomes
10:  for  $o = 1$  to  $O$  do
11:    INDUCETREE( $D_o$ )
12:  end for
13: end function

```

In order to overcome the issues of the fully induced tree that overfits the training data, as well as outliers and noise in the leaves of the tree, C4.5 introduced post-pruning. C4.5 replaces the tests in the deepest level of the tree with leaf nodes, where the prediction of the leaf node is the class that occurs the most in the new subset. A validation set is used to test the generalisation performance of unseen data and if the generalisation performance increases, the tree prunes away the decision node and replaces it with the leaf node. If the generalisation performance does not improve, the decision node is then re-inserted into the tree. This process is repeated until no more decision nodes are replaced. The tree is now robust to outliers and noise. *Data quality issues!*

D. Performance Metrics

Performance metrics are essential tools when the effectiveness of classification models are evaluated. Performance metrics provide a quantitative measure of how reliable and accurate a prediction model performs classification on a dataset. Key metrics include accuracy, precision, recall, and F1-score, each offering unique insights into different aspects of model performance [2].

a) *Accuracy*: Accuracy is a common method used to evaluate the performance of classification models. The accuracy of a predictive classification model is determined by the proportion of correctly predicted labels against the total number of predictions. The calculation of the accuracy of a predictive model is as follows:

$$\text{Accuracy} = \frac{\text{Number of Correct Predictions}}{\text{Total Number of Predictions}} \quad (21)$$

Accuracy is a popular choice of performance measure mainly because it is fairly easy to understand and compute. Accuracy generally perform well on well balanced datasets. On imbalanced datasets, accuracy can produce values that are misleading. If a model predicts

b) *Precision and Recall*: Precision is the proportion of true positive (TP) predictions against all of the TP and false

positive (FP). The equation to calculate the precision of a classification model is as follows:

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}} \quad (22)$$

Recall is the proportion of TP predictions against all of the TP and false negative (FN). The equation to calculate the recall of a classification model is as follows:

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}} \quad (23)$$

c) *F1-score*: The F1-score, also known as the Dice similarity coefficient, is the harmonic mean of precision and recall, that provides a balance between the precision and recall [4]. In multiclass classification precision and recall are adapted from binary classification to handle multiple classes. Instead of focusing on a single positive and negative class, these metrics are calculated for each class, where each class is treated as the positive class while considering all others as negative. The F1-score is calculated for each class and then averaged. There are three types of averaging F1-score metrics.

The first F1-score metric is the macro averaging F1-score, where the F1-score is calculated independently for each label and then these scores are averaged, giving each class the same weight. The equation used to calculate the macro F1-score is as follow:

$$F1_{macro} = \frac{1}{C} \sum_{i=1}^C \frac{2 \times \text{Precision}_i \times \text{Recall}_i}{\text{Precision}_i + \text{Recall}_i} \quad (24)$$

where C is the number of classes.

The second F1-score metric used is the micro averaging F1-score, where the contributions of all of the classes are aggregated inore to compute the total true positives, false positives and false negatives. The equation used to calculate the micro F1-score is as follow:

$$F1_{micro} = \frac{2 \times \text{Precision}_{micro} \times \text{Recall}_{micro}}{\text{Precision}_{micro} + \text{Recall}_{micro}} \quad (25)$$

Micro averaging treats each instance equally and is therefore useful when a dataset contains imbalanced classes.

The last F1-score metric is known as the weighted averaging F1-score. The weighted F1-score computes the F1-score for each class, just like in the macro F1-score. Each score is then weigh by the number of true instances in each class. The equation used to calculate the weighted F1-score is as follow:

$$F1_{weighted} = \sum_{i=1}^C w_i \frac{2 \times \text{Precision}_i \times \text{Recall}_i}{\text{Precision}_i + \text{Recall}_i} \quad (26)$$

where w_i is the proportion of true instances for label i . The equation to calculate w_i is as follows:

$$w_i = \frac{TP_i + FN_i}{\sum_{j=1}^C (TP_j + FN_j)} \quad (27)$$

where the denominator is the sum of all the true positives and false negatives of all classes, TP_i is the TP for class i and FN_i is the FN for class i .

When the distribution of the classes in a dataset is imbalanced, the F1-score is considered a more optimal performance metric than the accuracy of a model given that the weighted F1-score, macro F1-score and micro F1-score considers both false positives and false negatives.

III. METHODOLOGY

This section provides the approach taken to explore and analyse the dry beans dataset to identify and address data quality issues. Additionally, it details the methodology used to classify dry bean types through the use of a KNN and classification tree model.

A. Data Overview

The structure and characteristics of the Dry Beans dataset are thoroughly investigated. The details on the number of instances and types of features are investigated to provide a foundational understanding for the subsequent data analysis.

The Dry Bean dataset contains 13611 instances and 22 features. Among these features, 19 are numerical features and 3 are categorical features, with the 'Class' feature being the target feature. A closer examination reveals that missing values in the dataset are represented by '?'. Missing values identified by '?' is very problematic as it changes the continuous features containing missing values into discrete features. The dataset is modified so that each instance of a missing value is changed from '?' to a value of not a number (NaN). The next step is to identify some important characteristics of each feature and investigate these characteristics. The important characteristics of each continuous feature are summarised in Table II.

The important characteristics of each discrete feature in the Dry Bean dataset is summarised in Table III.




TABLE II: Continuous Features

Features	Count	% Miss.	Card.	Min.	1 st Qrt.	Mean	Median	3 rd Qrt.	Max.	Std Dev.
Area	13611	0.0	12011	20420	36328	53048.2845	44652	61332	254616	29324.0958
Perimeter	13611	0.0	13351	524.736	703.5235	855.2835	794.941	977.213	1985.37	214.2897
MajorAxisLength	13611	0.0	13543	183.601165	253.303633	320.1419	296.883367	376.495012	738.860153	85.6942
MinorAxisLength	13611	0.0	13543	122.512653	175.848170	202.2707	192.431733	217.031741	460.198497	44.9701
AspectRatio	13611	0.0	13543	1.024868	1.432307	1.5832	1.551124	1.707109	2.430306	0.2467
Eccentricity	13611	0.0	13543	0.218951	0.715928	0.7509	0.764441	0.810466	0.911423	0.0920
ConvexArea	13611	0.0	12066	-30	36714.5	53765.6926	45178	62294	263261	29778.0094
EquivDiameter	13611	0.0	12012	0.161417	215.068003	476.2541	238.438026	279.452162	3014441	25836.8656
Extent	13611	0.044082	13529	0.555315	0.718641	0.7497	0.759874	0.786852	0.866195	0.0491
Solidity	13611	0.0	13526	0.919246	0.985670	0.9871	0.988283	0.990013	0.994677	0.0047
Roundness	13611	0.0	13543	0.489618	0.832096	0.8733	0.883157	0.916869	0.990685	0.0595
Compactness	13611	0.132246	13525	0.640577	0.762577	0.7999	0.801291	0.834270	0.987303	0.0617
ShapeFactor1	13611	0.0	13543	0.002778	0.005900	0.0066	0.006645	0.007271	0.010451	0.0011
ShapeFactor2	13611	0.0	13543	0.000564	0.001154	0.0017	0.001694	0.002170	0.003665	0.0006
ShapeFactor3	13611	0.0	13543	0.410339	0.581359	0.6436	0.642044	0.696006	0.974767	0.0990
ShapeFactor4	13611	0.0	13611	0.695579	1.614151	2.3681	2.368757	3.115695	3.966119	0.8716
ShapeFactor5	13611	0.0	13543	0.947687	0.993703	0.9951	0.996386	0.997883	0.999733	0.0044
ShapeFactor6	13611	0.036735	13606	0.000466	45.258826	89.3586	88.76667	134.273148	178.985023	51.8386
Sort order	13611	0.0	13611	0.000089	0.248187	0.5003	0.50381	0.750096	0.999985	0.2879

TABLE III: Categorical Features

Features	Count	% Miss.	Card.	Mode	Mode Freq.	Mode %	2 nd Mode	2 nd Mode Freq.	2 nd Mode %
Constantness	13611	0.0	2	1	12289	90.287268	0	1322	9.712732
Colour	13611	0.044082	4	brown	6115	44.926897	black	3541	26.015723
Class	13611	0.124899	7	DERMASON	3542	26.02307	SIRA	2634	19.351995

The most important observation made from Table III is that the 'Class' feature is imbalanced, as the first and second modes together form 45% of the data. The five other classes makes up the remaining 55%, and thus the 'Class' feature is imbalanced.

Several points of interest emerge after a closer examination of Table II. The first observation made is that the feature 'Sort Order' feature has a cardinality of 13611, a minimum of approximately 0, a first quartile value of approximately 0.25, a mean of approximately 0.5, a third quartile of approximately 0.75 and a maximum value of approximately 1. These statistics all indicate that 'Sort Order' is a unique identifier or index. As a result, 'Sort Order' may be considered as an irrelevant feature for predictive modeling purpose.

The second observation made is that feature 'ConvexArea' contains a negative number as the minimum value. This is impossible as area can not be negative. The boxplot of feature 'ConvexArea' is given by Figure 1 to further investigate the properties of the feature 'ConvexArea'.

Figure 1 clearly shows that the negative value is an obscure outlier, thus can this instance of the dataset be classified as an invalid outlier and should either be imputed or dropped from the dataset.

The third observation made is that both the minimum and the maximum values of feature 'EquivDiameter' deviates largely from the rest of the statistics. This large deviation may indicate that these instances of the dataset are also invalid outliers. Figure 2 is the boxplot of feature 'EquivDiameter' that helps to investigate the obscure values of the instances in the dataset.




Fig. 1: Boxplot of feature ConvexArea vs feature Class




Fig. 2: Boxplot of feature EquivDiameter vs feature Class

From Figure 2 it is quite clear that the maximum value of feature 'EquivDiameter' is an invalid outlier. This instance is dropped from the dataset and the boxplot is generated again, as shown in Figure 3, to investigate the minimum value of the feature 'EquivDiameter'.




Fig. 3: Boxplot of feature EquivDiameter vs feature Class after the maximum value was dropped

As seen in Figure 3 there is another instance in the dataset that deviates significantly from the rest of the observations and can also be classified as an invalid outlier. This outlier is also removed from the dataset and the boxplot is again generated, as shown in Figure 4, to investigate the minimum value of the feature 'EquivDiameter'.




Fig. 4: Boxplot of feature EquivDiameter vs feature Class after the invalid outliers were dropped

From Figure 4 it is clear that the minimum value is also an invalid outlier that should either be imputed or removed from the dataset.

B. K-Nearest Neighbours

a) *Data preparation:* Data preparation is a very crucial step that must be completed before the KNN algorithm can be applied to the Dry Bean dataset. If a KNN model is implemented as described in Section II-B, then will the KNN model be robust to missing values, noise and outliers in the

dataset. To prepare the Dry Bean dataset for a KNN model, the missing values have to be replaced with NaN and the instances that contains missing values in the 'Class' feature is dropped.

KNN is an instance-based learning algorithm that calculates the distance between feature instances to classify new data points based on the majority class of their nearest neighbours. Therefore, all features need to be numeric for the distance calculations to be valid. The 'Colour' feature is a categorical feature with no natural ordering and will therefore be one hot encoded.

This implementation of KNN is sensitive to imbalanced classes. The 'Class' feature is imbalanced, as discussed in Section III-A. Tomek links and SMOTE resampling techniques are applied to the dataset in order to address this imbalance and improve the performance of the KNN classification model. Additionally, standardisation is applied to scale the feature values, to ensure that all features contribute equally to the distance calculations and to enhance the accuracy of the model. Standardisation is applied by the use of Equation (6).

Subset feature selection has to be applied to the newly resampled and standardised dataset. If a KNN algorithm is fitted with all of the input features, it becomes very likely that the algorithm overfits, thus having a poor generalisation on unseen data. In Section II-A2b, feature selection is discussed in detail. By the use of a wrapper feature selection technique, the dimension of the Dry Bean dataset can be reduced to the construction of a KNN algorithm that overfits the dataset.

b) *KNN model implementation:* The KNN model for predictive analysis on the Dry Bean dataset is implemented as described in Section II-B. An additional integer control parameter has been introduced to Algorithm 1 that allows for the selection of the appropriate distance metric to be applied in the process of model execution. The implementation of the KNN model is represent by Algorithm 3.

Algorithm 3 Improved k-Nearest Neighbours (kNN)

```

1: function KNN( $D, \mathbf{x}, k, p$ )
2:   for all  $\mathbf{x}_i \in D$  do
3:      $\mathbf{d} = \text{DISTANCE}(\mathbf{x}_i, \mathbf{x}, p)$ 
4:   end for
5:    $\text{SORT}(\mathbf{d})$ 
6:    $S = \text{set of } k \text{ patterns in } D \text{ closest to } \mathbf{x}$ 
7:   Return class as majority class in  $S$ 
8: end function

```

where the equation to calculate the distance between two input features of instances \mathbf{x}_1 and \mathbf{x}_2 is as follows:

$$d(\mathbf{x}_1, \mathbf{x}_2) = \left(\sum_{n=1}^N |x_{1n} - x_{2n}|^p \right)^{\frac{1}{p}} \quad (28)$$

The Equation (28) is observed to be the Minkowski distance metric. If p is equal to one, the Manhattan distance is calculated between \mathbf{x}_1 and \mathbf{x}_2 . If p is equal to two, the Euclidean distance is calculated between \mathbf{x}_1 and \mathbf{x}_2 and if p is larger than 2, the Minkowski distance is calculated between \mathbf{x}_1 and \mathbf{x}_2 .

assuming numeric distance calculation.

skew class distribution?

35

depends on p.
why?
Do you state a question or a fact?

C. Classification Tree

a) *Data preparation*: Data preparation is a crucial step that must be completed before the classification tree algorithm can be applied to the Dry Bean dataset. If a classification tree algorithm is implemented as described in II-C, then will the classification tree be robust to outliers, missing values and noise in the dataset. To prepare the Dry Bean dataset for a classification tree model, the missing values have to be replaced with NaN and the instances that contains missing values in the 'Class' feature is dropped. The classification tree is sensitive to class imbalances. Instead of resampling the data, another technique is implemented within the classification tree algorithm to enhance the algorithm's robustness to class imbalances. *Why not self-label? from less splitting for tree?*

b) *Classification Tree model implementation*: The classification tree model for predictive analysis on the Dry Bean dataset is implemented as described in Section II-C. The inclusion of these two additional control parameters, being the maximum depth of the tree and the minimum number of samples needed to create a split in the tree, creates a tree that is less prone to overfitting. In Algorithm 4 the two additional control parameters are added.

Algorithm 4 Improved Classification Tree Induction

```

1: function INDUCETREE( $D$ ,  $depth$ )
2:   if  $|D| = 0$  then
3:     Return leaf with default class
4:   end if
5:   if  $|D| < \text{min\_sample\_split}$  or  $depth \geq \text{maxdepth}$ 
   then
6:     Return leaf with class label  $y_m$ , containing  $D$ 
7:   end if
8:   if  $|D| > 0$  and  $\forall \mathbf{x} \in D$  the class is the same, i.e.  $y_m$ 
   then
9:     Return leaf with class label  $y_m$ , containing  $D$ 
10:  end if
11:  Select a test based on a single input variable
12:  Split  $D$  into  $D_1, D_2, \dots, D_o$ , where  $O$  is the number of
   outcomes
13:  for  $o = 1$  to  $O$  do
14:    INDUCETREE( $D_o$ ,  $depth + 1$ )
15:  end for
16: end function

```

An additional return condition was added to reduce the complexity of the model and to create a classification tree that is less prone to overfit. If the current depth $depth$ is greater than or equal to the maximum depth control parameter, or if the sample size is smaller than the minimum sample split, the tree returns the majority class in the data.

The classification tree makes use of pruning by separating the training set into a training and validation set. The tree is induced by the use of training set and then pruned by use of the validation set. Algorithm 5 explains the algorithm used to prune the tree.

Algorithm 5 Simplified Post-Pruning of a Decision Tree

```

1: function PRUNETREE( $X_{val}$ ,  $y_{val}$ )
2:   if Tree is not trained then
3:     raise ValueError("The tree has not been trained yet.")
4:   end if
5:   PRUNENODE(root,  $X_{val}$ ,  $y_{val}$ )
6: end function
7: function PRUNENODE(node,  $X_{val}$ ,  $y_{val}$ )
8:   if node is a leaf then return
9:   end if
10:  for each branch in node.branches do
11:    Determine threshold and mask for  $X_{val}$  based on
    the branch value.
12:    Subset  $X_{val}$  and  $y_{val}$  using the mask.
13:    if the subset is not empty then
14:      PRUNENODE(branch,  $X_{val\_subset}$ ,  $y_{val\_subset}$ )
15:    end if
16:  end for
17:  if  $X_{val}$  is not empty then
18:    Compute original predictions and original F1-
    score.
19:    Temporarily make the node a leaf and set the
    prediction to the majority class.
20:    Compute pruned predictions and pruned F1-score.
21:    if pruned F1-score < original F1-score then
22:      Revert node to its original state.
23:    end if
24:  end if
25: end function

```

F1-score is used as the performance metric to compare the predictions of the original tree to the predictions of the pruned tree. If the F1-score of the pruned tree is lower than the F1-score of the original tree, the changes are reverted to ensure that pruning does not degrade the performance of the model.

IV. EMPIRICAL PROCEDURE

A. Performance Metrics

Accuracy is used to evaluate the performance of the KNN model, as the dataset has been resampled to have balanced class distributions. As for the Classification tree, the performance measures used to evaluate the model will be the weighted F1-score, since the dataset has an imbalanced class distribution. *8 ROC, AUC confusion matrix?*

B. Control Parameters

The KNN experiments are performed by the use of the euclidian distance metric and a value of 147 for k . The experiments of the classification tree are performed by the use of a max depth of 9 and a minimum sample split of 2.

C. Experimental Setup

To evaluate the performance of both the KNN and classification tree algorithms, a 5-fold cross validation technique was applied, that splitted the Dry Bean dataset into five equally

distributed datasets. Four of these datasets were used to train each model and the remaining dataset was used to test the classification algorithm. This ensured that the evaluation was fair and consistent across the entire dataset and the bias that might arise from relying on a single training-test split was reduced. The rotation of the test set, allowed for a more comprehensive understanding of the performance of the model on unseen data.

D. Statistical Significance and Analysis

The statistical significance of the results from the KNN and classification tree models is determined by the evaluation of the accuracy and weighted F1-score respectively of each of the models implemented by the use of a 5-fold cross-validation over a grid of different possible values that the control parameters could be. By use of the Kruskal-Wallis test to check for stochastic dominance amongst the samples. The paired t-test was then performed to obtain this optimal value of the control parameters for both the KNN and classification tree models. A significance level of 0.05 is used.

V. RESEARCH RESULTS

This section provides an overview of the experimental results collected and the outcomes of the empirical analysis conducted on the DryBeanDataSet by use of the implemented classification algorithms: KNN and decision tree.

A. Correlation Matrix

The values in this correlation matrix ranges from negative one to one, where a value of one indicates a strong positive linear relationship between the features, a value of negative one indicates a strong negative linear relationship between the features and a value of zero indicates that the features have no linear relationship. Figure 5 shows the correlation matrix of the features in the Dry Beans dataset after the 'Colour' feature has been one hot encoded and the 'Class' feature has been transformed into an ordinal encoded feature.

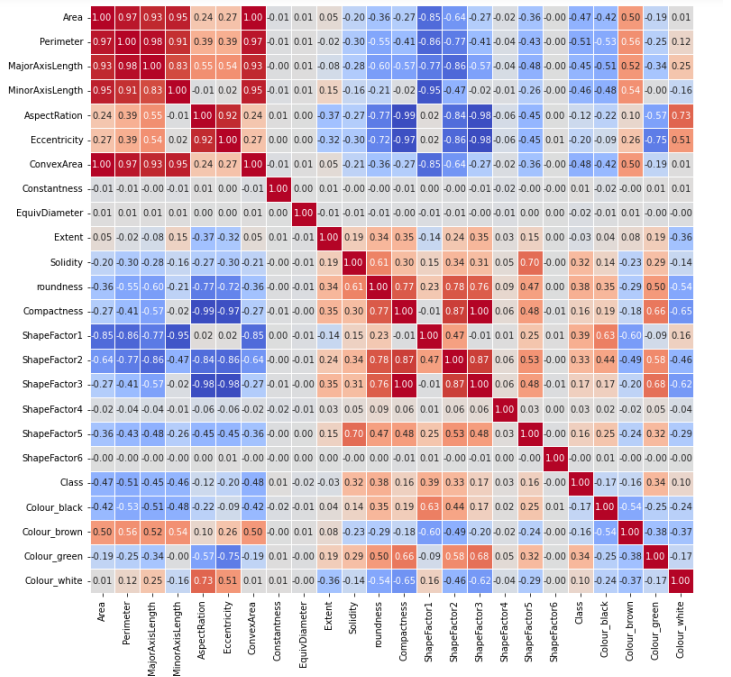


Fig. 5: Correlation matrix of the Dry Bean dataset features

After the examination of Figure 5, there are a few observations that can be made. The first observation is that the features 'Area', 'Perimeter' and 'MajorAxisLength' are strongly positive correlated features with one another and 'ShapeFactor1' and 'ShapeFactor2' are strongly negative correlated features with one another. The features 'Constantness' and 'Equivdiameter' have low or near-zero correlations with many features, suggesting that these features might be irrelevant features

B. Best Subset for KNN

When backwards subset selection is implemented, the model starts as a dataset that contains all of the features. One feature is dropped from this set and a KNN model is created from the new subset. The accuracy score is then calculated for this subset as the inner criteria and the feature is added back into the subset. Then the next feature is dropped and the process follows the same. This process is done until all of the features have been dropped and a KNN model has been implemented for each smaller subset. The subset that obtained the highest score is then considered the best subset of the features with one less feature than the original dataset has. This best feature subset is then taken and the process is repeated until there is only a subset left with one feature. This process of the inner criteria helps reduce the bias of the model.

The subsets obtained in the process above is then taken and an outer criteria is applied to each one of them. The outer criteria in this case is a k-fold cross-validation. The scores of each of these k-fold iterations were computed and the mean of the scores is calculated. The mean error is then calculated as well as the standard deviation of the mean error. In Figure 6 the mean cross-validation error is represented by the dot and the error bar is the standard deviation of the score errors.




Fig. 6: Cross validation error of the different sized subsets

When the one standard deviation rule is applied to the graph, it is seen that the smallest subset of features that falls within one standard deviation of the feature with the lowest score, in this case the subset with 18 features, is the subset with 14 features. Hence the best subset to use for the KNN is the subset with 14 features.

C. Hyperparameters

a) *K Nearest Neighbours*: The KNN model has two hyperparameters that needs to be tuned. The first of which is the value of p , that indicates which distance metric should be used and the second parameter is the value of the number of nearest neighbours or k . By first tuning the value of p using a grid search, it is found that the Euclidean distance performed better on the Dry Bean dataset. After obtaining the value of p , 50 KNN models were fitted to the grid search using a 5-fold cross-validation to determine the optimal value for k . The 50 values that k is tested on is the values between 85 to 185 with an increment of two between each value. The Figure 7 shows how each KNN model performed on average after a 5-fold cross-validation.




Fig. 7: Mean test scores vs the number of neighbours

Running a Kruskal-Wallis test on these 5-fold cross-validation scores yielded a Kruskal-Wallis H-statistic equal to 141.9776 and a p-value of $5.5326e - 11$, which indicates that there is a significant difference between the performance of

different k values. To identify where this significant difference lies, a paired t-test statistic is used. From a KNN model with 145 number of neighbours to a KNN model with 147 number of neighbours, the paired t-test returns a p-value of 0.0341. The paired t-test and the Kruskal-Wallis p-value is smaller than the significance level of 0.05 and therefore can the optimal value of k be chosen as 147.

b) *Classification Tree*: The two hyperparameters that needs to be tuned for the classification tree is the maximum depth of the tree and the minimum sample split. The dataset used on the classification tree is unbalanced and therefore, the weighted F1-score is used to perform the statistical tests. After all of the 5-fold cross-validation scores are obtained and the paired t-test has been computed for a grid search where the maximum depth was tested to be [8, 9, 10, 11, 12] and the minimum sample split is tested to be [2, 3, 4, 5], the most optimal value obtained for each of these grids is 9 for the maximum depth and 2 for the minimum sample split.

D. Model Evaluation

The final models to be evaluated is a KNN model with the Euclidean distance metric and a k value of 147. Aswell as the final classification tree model with a maximum depth of 9 and a minimum sample split of 2.

After inducing the tree and post-pruning is applied, the tree that is used to perform classification tasks on is given by Figure 8.

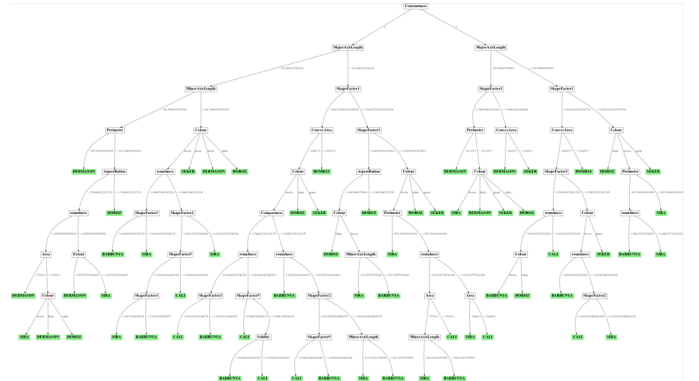


Fig. 8: Classification tree

In Table IV, all of the performance metrics are shown on how the KNN performed on the Dry Bean dataset.

TABLE IV: Classification report for the KNN

	precision	recall	f1-score	support
BARBUNYA	0.97	0.94	0.95	651
BOMBAY	1.00	1.00	1.00	639
CALI	0.95	0.97	0.96	634
DERMASON	1.00	0.99	1.00	629
HOROZ	1.00	1.00	1.00	655
SEKER	1.00	0.99	1.00	635
SIRA	0.97	0.99	0.98	643
accuracy			0.98	4486
macro avg	0.98	0.98	0.98	4486
weighted avg	0.98	0.98	0.98	4486

In Table V, all of the performance metrics are shown on how the KNN performed on the Dry Bean dataset.

TABLE V: Classification report for the classification tree

	precision	recall	f1-score	support
BARBUNYA	0.86	0.91	0.88	289
BOMBAY	1.00	1.00	1.00	98
CALI	0.93	0.89	0.91	323
DERMASON	0.95	0.98	0.97	698
HOROZ	0.99	0.99	0.99	383
SEKER	0.99	0.99	0.99	391
SIRA	0.97	0.92	0.95	536
accuracy			0.96	2718
macro avg	0.96	0.96	0.96	2718
weighted avg	0.96	0.96	0.96	2718

Figure 9 shows the classification matrix of the KNN model obtained after classification of the test set.




Fig. 9: Confusion matrix obtained by the KNN model

Figure 10 shows the classification matrix of the classification tree model obtained after classification of the test set.




Fig. 10: Confusion matrix obtained by the KNN model

After the evaluation of the models, it is clear that KNN is the better prediction model for the Dry Bean dataset. This is to be expected, considering that most of the features in the Dry Bean dataset is numerical features and that the classification tree groups all unique values of the numerical features into

categorical features. KNN calculates the distance between the instances of each feature, so it makes sense why KNN performs the best on the Dry Bean dataset.

VI. CONCLUSION

In this report, the Dry Bean dataset was thoroughly analysed and the important characteristics of each feature was investigated to better understand the data and to identify any data quality issues. The Dry Bean dataset was used to construct two classification models, namely a k-nearest neighbours (KNN) and a decision tree. Before the KNN model was trained on the dataset, some data quality issues had to be addressed such as resampling by use of Tomek links and synthetic minority oversampling technique (SMOTE) to balance the class distributions, standardisation of the numerical input features and best subset selection using backwards subset selection. The KNN algorithm implemented is robust to missing values, outliers and noise in the dataset. The decision tree is induced to overfit the data and separate the classes with 100% accuracy and by use of post-pruning, to optimise the tree and make the model less prone to overfit.

By the use of k-fold cross-validation and performance metrics such as accuracy for the KNN model, because of the evenly distributed classes and F1-score for the decision tree classification model, because of the imbalanced distributed classes, it was shown that both classification models can effectively be applied to the Dry Bean dataset. After the implementation of a grid search and the appropriate statistical tests, to determine the optimal values for the number of neighbours and the optimal distance function to use, the KNN model produced a number of neighbours equal to 147 and the distance metric used was the Euclidean distance metric. The optimal values for the maximum depth and minimum sample split for the decision tree after the implementation of a grid search and the appropriate statistical tests came out to be 9 and 2 respectively.

In conclusion, both the KNN and decision tree models are viable methods for the task of classification on the Dry Bean dataset. However, the model that performs better classifications on the Dry Bean dataset is the KNN classification model.

¹

REFERENCES

- [1] A. C. Acock "Working with missing values." In: Journal of Marriage and family (2005).
- [2] J. Braet, M. Cristina, Hinojosa-Lee, and J. Springael. "Evaluating performance metrics in emotion lexicon distillation: a focus on F1 scores." (2024).
- [3] L. Breiman, J. Friedman, R. A. Olshen and C. J. Stone "Classification and Regression Trees." In: Routledge. (1984).
- [4] B. J. Erickson, and K. Felipe "Magician's corner: 9. Performance metrics for machine learning models." In: Radiology: Artificial Intelligence 3, no. 3 (2021): e200126.
- [5] M. C. Jones, B. W. Silverman "E. Fix and J.L. Hodges (1951): An Important Contribution to Nonparametric Discriminant Analysis and Density Estimation: Commentary on Fix and Hodges (1951)." In: International Statistical Review / Revue Internationale de Statistique, vol. 57, no. 3, (1989).

¹Code to the project can be found here: <https://github.com/Andrevan071/ML441Assignments>

- [6] J. R. Quinlan "Induction of decision trees" In: Mach Learn 1 (1986).
- [7] J. R. Quinlan "C4. 5: programs for machine learning." In: Elsevier, (1993).
- [8] J. W. Tukey. "Exploratory data analysis." In: Reading/Addison-Wesley (1977).